ConfiGuard-Lite: Repository Cleanup Report

Date: 2026-11-11 | Branch: main @ 7c1fdd7 | Space freed: ~24.1 MB

Summary

I checked the whole repository: all 372 tracked files, every ignored or untracked path, and the code that refers to each candidate folder. Removed: **only** generated caches, stale build output and empty folders. Each one either gets rebuilt automatically or can be regenerated by an existing script. **No tracked (committed) file was deleted or changed.** `git status` was the same before and after the cleanup. Your one uncommitted change, frontend/package-lock.json, was left alone.

What was removed and why

Path	Size	Why it was safe to remove
34 x __pycache__/ folders (src/, tests/, scripts/; .venv left alone)	~4 MB	Python bytecode caches (.pyc). The interpreter rebuilds them automatically the next time the code runs. Already git-ignored, never committed.
.pytest_cache/	71 KB	pytest's last-failed/step-wise state. It is rebuilt on the next test run. Git-ignored.
frontend/test-results/	1 KB	Leftover Playwright output (only .last-run.json). Recreated on each e2e run. Git-ignored.
build/hf-space/ (221 files)	24 MB	Generated copy of the Hugging Face Space bundle, dated 22 Oct. It is **stale**: later commits (the FFmpeg error reporting and 4K face-detection fix) are not in it, so pushing it would ship old code. It has no .git folder, so no deploy history was lost. Regenerate with `python scripts/prepare_hf_space.py` (output defaults to build/hf-space).
.github/workflows/ (empty)	0 B	Empty folder with no workflow files. Git does not track empty folders, so it did nothing.

What was deliberately KEPT (and why)

Path	Size	Why it must stay
build/release/	20 MB	**Critical.** checkpoints/ is empty, so this is the **only local copy of the trained student model** (student_p80 ONNX), its calibration, quality gate, YuNet face detector and C2PA trust list. The HF Space and Oracle bundles are built from it.
outputs/onnx/	25 MB	EfficientNet-B0 / MobileNetV4 baseline ONNX exports. Benchmarks and comparisons use them, and re-creating them means downloading the models again.
models/face_detection/	228 KB	YuNet face detector. The media pipeline needs it.
frontend/dist/	2.9 MB	Production React build. The local FastAPI server serves it automatically. Deleting it would quietly switch the app to the older fallback UI until you rebuild it.
frontend/node_modules/	487 MB	Can be regenerated with npm ci, but it is needed to build or test the frontend. Delete it yourself only if disk space is tight.
.venv/	5.6 GB	Your Python environment (PyTorch/CUDA). Required.
.env	4 KB	Your local secrets and paths. Never touch.
.agents/, skills-lock.json	25 KB	A frontend design skill you installed for AI agents. It is your tooling, not project clutter.
quality v1 / v2 / hybrid modules + scripts	-	Committed research code with tests. DECISIONS.md and EXPERIMENT_LOG.md refer to it, and it is needed to reproduce those results. The service uses v1.

Path	Size	Why it must stay
service/static/ (old UI)	-	The fallback UI the server uses when frontend/dist is missing.
PROJECT_OVERVIEW.md + .pdf	180 KB	Both added on purpose in the last two commits.
cache/, checkpoints/, data/, models/, outputs/ .gitkeep	0 B	Placeholders that keep the folder layout the code expects.

Verification

After the cleanup I ran the full Python test suite, with bytecode and pytest caches turned off so the deleted caches would not come straight back:

`PYTHONDONTWRITEBYTECODE=1 .venv/Scripts/python.exe -m pytest -q -p no:cacheprovider -x`

Result: 715 passed in 236 s, 0 failures.

git status before = after: only `M frontend/package-lock.json` (your existing change).

Recommendations

1. **Back up build/release/** (for example to cloud storage or an HF model repo). It is the only copy of your trained model and is not in git.
2. Run `scripts/prepare_hf_space.py` again before your next Hugging Face deploy, so the Space includes the latest fixes.
3. Optional: delete frontend/node_modules (487 MB) if you need the space, and restore it later with `npm ci`.